

4DW-VLA: 4D and World-Model Aligned Vision-Language-Action for Robotic Manipulation

Gang Luo, *Member, IEEE*, and Daohe Tongtai

Abstract—Most vision-language-action (VLA) policies operate on 2D images, yet robots act in a metric 3D workspace whose configuration evolves over time, a 4D problem. Methods that inject 3D or 4D structure often rely on depth sensors, point-cloud encoders, or heavy scene reconstruction, which increases model size, training cost, and deployment latency; many of them still treat spatial geometry and temporal dynamics as separate concerns. World-model-augmented VLAs predict future 2D video but rarely align that visual foresight with the robot’s own metric body motion, so predicted actions can be temporally coherent in pixel space yet physically inconsistent in 3D. We present 4DW-VLA, a 4D and world-model aligned VLA whose central component is a dedicated 4D Future Expert that predicts future 7D link trajectories (3D position + orientation) from forward kinematics (FK) at negligible cost, orders of magnitude cheaper than point-cloud or depth-pipeline 4D extraction. Unlike prior FK-based approaches that discard the 4D branch after training, the 4D Future Expert is retained at deployment and continues to supply metric body-motion context to the action head. A compact World Knowledge Expert exchanges information with the 4D expert through Unified Bidirectional Cross-Attention (UBCA), producing 2D-foresight tokens aligned with predicted 4D motion; these tokens supervise a frozen large video world model during training only. At inference, the world model is removed; the VLM, the two lightweight experts, and the action head remain, yielding a small deployable policy with no video-generation cost. On three RoboTwin 2.0 bimanual tasks, 4DW-VLA matches or exceeds representative baselines from the 2D-foresight and depth-plus-video VLA families, with the largest gains under domain randomization.

Index Terms—Vision-language-action models, 4D representation, world models, robotic manipulation, cross-attention.

I. INTRODUCTION

LEARNING-BASED manipulation policies and VLA models [2, 3, 5, 29, 11, 1] have improved language following and cross-task generalization, but they still under-represent the physical world in which a robot moves.

Most VLAs are fundamentally 2D. They consume RGB images and language, yet manipulation unfolds in a 3D workspace; when time is included, the relevant object is 4D structure: how 3D geometry changes across the action horizon. A policy that reasons only in image space lacks explicit metric body motion and cannot guarantee that predicted actions respect the robot’s own spatial evolution.

3D-/4D-aware VLAs often pay a heavy geometric tax. Point clouds, depth maps, voxel grids, and 3D Gaussian fields can improve spatial reasoning [27, 26, 6], but they typically

require extra sensors or large perception networks to extract and encode geometry at every timestep. The resulting models are bigger, slower to train, and harder to deploy. Moreover, many 3D-enhanced designs treat geometry as a **static** per-frame cue and under-model **temporal** changes in the robot’s own configuration.

Temporal prediction is usually misaligned across modalities. Actions are inherently future-oriented: a policy outputs action chunks over horizon H . Yet most VLAs do not jointly predict how **2D observations** and **3D/4D body state** should co-evolve over that horizon. World-model and world-action-model (WAM) lines predict future video [7, 14, 18], but often without explicit metric 4D body motion; the camera may “imagine” a future that is not consistent with the robot’s actual 3D trajectory (e.g., the arm rises in joint space but the predicted video does not show the raised hand).

These signals should be coupled. When a robot lifts its end-effector in 3D, the egocentric camera stream should reflect that motion; conversely, visual foresight should not contradict feasible body dynamics. Effective manipulation therefore requires **aligning** (a) future **4D body motion** in the robot frame and (b) future **2D video** the camera will see, and using both to predict (c) future **actions**.

We propose **4DW-VLA**, a **4D and World-model aligned VLA**, to address the above gaps. Our design rests on three ideas:

- 1) **A dedicated 4D Future Expert, retained at deployment.** Instead of reconstructing scene geometry from depth or point clouds, we compute 7D link trajectories (3D position + orientation) directly from proprioception via FK, at orders-of-magnitude lower cost than perception-based 4D pipelines. Unlike prior FK-based approaches that treat 4D as a training-time auxiliary discarded at inference, the 4D Future Expert is a first-class attention path that **remains active at deployment**, continuously supplying metric body-motion structure to the action head.
- 2) **4D–2D alignment via bidirectional cross-attention.** The 4D Future Expert and a compact **World Knowledge Expert** exchange information through **Unified Bidirectional Cross-Attention (UBCA)**: at every timestep, each expert attends causally within its own stream and to past tokens of the other, coupling 4D body foresight and 2D visual foresight before action decoding.
- 3) **Train with a large frozen world model; deploy without it.** The World Knowledge Expert produces compact

implicit 2D-foresight tokens already aligned with 4D predictions. During training, these tokens condition a **frozen** large video diffusion world model; at inference the world model is **discarded**, so the deployable VLA stays small and fast while retaining distilled, 4D-aligned visual dynamics.

Fig. 1 summarizes the architecture. We co-train four objectives (current 3D state estimation, future 4D prediction, future 2D foresight distillation, and action-chunk generation) under one weighted multi-task loss (Section III-E). The video world model g_{wm} stays **frozen**; the VLM backbone is fine-tuned at **one-fifth** the learning rate of the expert and action heads over **10,000** optimizer steps.

Contributions:

- 1) We introduce a **4D Future Expert** that uses **FK-derived 7D link trajectories** as a lightweight body-centric 4D representation and is **retained at deployment**, unlike prior FK-based approaches that discard the 4D branch after training (Section III-B).
- 2) We couple the 4D Future Expert with a compact **World Knowledge Expert** through **UBCA**, aligning future metric body trajectories with future egocentric video to jointly condition chunked action prediction (Section III-C).
- 3) We distill a **frozen** video world model into trainable foresight tokens and **remove** it at inference, preserving 4D–2D alignment without deployment-time video generation (Sections III-C–III-F).

II. RELATED WORK

A. Foundation Models in Robotics

Generalist robot policies are trained on large multi-task datasets [25, 10, 24] and increasingly built on vision-language priors [3, 11, 1, 13]. RT-2 [3] and OpenVLA [11] cast actions as language-model tokens; Diffusion Policy [5], RDT [13], and π_0 [1] generate continuous action chunks via diffusion or flow matching. These foundations excel at semantic generalization but largely treat observations as 2D images without explicit 4D body–video alignment.

World-action models (WAMs) co-train video prediction with control [14, 18, 16]: the predicted future regularizes the action pathway, and generation is often dropped at test time. **4DW-VLA** follows the “train with foresight, infer without the generator” principle, but differs in **what** is aligned: we explicitly couple **FK-derived 4D body trajectories** with **2D video foresight** through bidirectional cross-attention before action decoding, rather than relying on video prediction alone.

B. 3D and 4D Representations for Manipulation

A growing line injects 3D or 4D structure into manipulation policies and VLAs. These methods improve spatial and temporal reasoning, but the **4D signals they use are typically heavy**: they require extra sensing, large encoders, auxiliary foundation models, or scene-scale reconstruction, which inflates compute, memory, and model size. **4DW-VLA** deliberately avoids this tax by deriving 4D body motion from FK and distilling 2D

dynamics into a small expert rather than embedding a full world model in the deployable policy.

RynnWorld-4D [28] builds a tri-branch diffusion world model that co-generates future RGB, **depth**, and **optical flow** from RGB-D input, curates hundreds of millions of frames with pseudo depth/flow labels, and trains an inverse-dynamics policy on internal 4D latents, an **infrastructure-heavy** approach. **4D-VLA** [22] injects sequential **RGB-D** frames with 3D coordinate embeddings into VLA pretraining; **depth must be available or estimated for every frame**. **Pri4R** [20] supervises VLMs with **privileged scene point tracks** from depth-based tracking, **expensive to obtain and encode**. **Lift3D-VLA** [19] lifts VLAs with explicit **point clouds**, GC-MAE pretraining over reconstructed geometry, and layer-wise temporal action modeling, adding **substantial parameters and preprocessing**. **SpatialVLA** [15] integrates **monocular depth estimation** into a 4B-scale VLA, introducing a **permanent depth branch** at inference unless distilled away. **4D-WAM** [17] aligns WAMs with dense **pixel-level 3D trajectory fields** from a large 4D foundation teacher.

In contrast, **4DW-VLA** obtains 4D through **FK on proprioception** (microseconds on CPU, no depth network, no point cloud) and aligns it to 2D foresight through **two small coupled experts**, using a frozen world model only as a **training-time teacher**. The deployable model excludes the teacher entirely.

C. World Models and World-Action Models

Classical world models predict future states for planning or representation learning [7, 12, 8]. Recent WAMs make the future **action-facing** [14, 21]: video or latent subgoals regularize the policy [18, 16]. **4DW-VLA** contributes a complementary axis: the training signal is not only “predict future pixels” but **align future pixels with future FK 4D body motion** via bidirectional cross-attention, then condition actions on both distilled pathways. The large video model is frozen and removed at deployment, similar in spirit to fast WAM recipes but with explicit **4D–2D coupling**.

III. METHOD

We present 4DW-VLA, a 4D and world-model aligned VLA for language-conditioned manipulation. A VLM encodes multi-view RGB and instructions; three trainable heads (a **4D Future Expert**, a **World Knowledge Expert**, and an **Action Expert**) cooperate through UBCA. Fig. 1 shows the overview.

A. Problem Formulation

Let o_t denote multi-view RGB, ℓ_t the language instruction, and q_t proprioception. Write p_t for the current metric 3D body state (FK-derived link poses), $P_{t+1:t+H}$ for the future 4D body trajectory, and $V_{t+1:t+T}$ for future egocentric video.

4D–2D future alignment:

$$p(P_{t+1:t+H}, V_{t+1:t+T} \mid o_t, \ell_t, q_t). \quad (1)$$

Action prediction:

$$p(A_t \mid P_{t+1:t+H}, V_{t+1:t+T}, o_t, \ell_t, q_t), \\ A_t = [a_t, \dots, a_{t+H-1}]. \quad (2)$$




Fig. 1: **4DW-VLA overview.** A VLM prefix encodes RGB and language. The **4D Future Expert** predicts current and future FK 7D link trajectories from proprioception and 4D history. The **World Knowledge Expert** predicts implicit 2D-foresight tokens aligned with 4D through **Unified Bidirectional Cross-Attention**. A **frozen** video world model supervises 2D foresight during training only. The **Action Expert** consumes VLM tokens, both experts' outputs, and predicts action chunks. At inference, the world model is dropped; the VLM, two experts, and action head remain.

At deployment we use compact representations Z_t^{4D} and Z_t^{wk} :

$$p(A_t | Z_t^{4D}, Z_t^{wk}, o_t, \ell_t, q_t). \quad (3)$$

Current 3D state estimation (auxiliary):

$$p(p_t | o_t, \ell_t, P_t^{hist}), \quad (4)$$

where P_t^{hist} is a short history of FK 4D tracks.

B. 4D Future Expert

Motivation. Point-cloud, depth, and scene-track 4D representations require sensing or heavy networks (Section II). Manipulators already expose q_t ; FK maps joints to link poses in microseconds. We therefore build 4D body supervision from **FK only**, gaining orders-of-magnitude lower latency and zero extra sensors compared with perception-based 4D pipelines.

7D link trajectories. For each link/keypoint $k \in \{1, \dots, K\}$,

$$x_t^k = \text{FK}_k(q_t) = [p_t^k, r_t^k] \in \mathbb{R}^7, \quad (5)$$

where $p_t^k \in \mathbb{R}^3$ is position and $r_t^k \in \mathbb{R}^4$ orientation (unit quaternion) in the robot base frame.

Expert architecture. The 4D Future Expert is a lightweight transformer over temporal 4D tokens derived from q_t and P_t^{hist} , plus conditioning from the VLM prefix. It outputs reconstructed **current** 4D state $\hat{P}_t$, predicted **future** trajectory $\hat{P}_{t+1:t+H}$, and hidden tokens Z_t^{4D} consumed by the World Knowledge Expert and Action Expert.

Role of q_t . Proprioception enters both as raw input tokens and as the quantity that defines FK targets. Conditioning on q_t anchors predictions in the robot's instantaneous configuration; through UBCA (Section III-C), Z_t^{4D} also supplies metric body context to the 2D foresight pathway.

4D future loss:

$$\mathcal{L}_{4D} = \beta \left(\text{MSE}(\hat{P}_t(q_t), P_t) + \gamma \text{MSE}(\hat{P}_{t+1:t+H}(q_t), P_{t+1:t+H}) \right). \quad (6)$$

Current state loss:

$$\mathcal{L}_{cur} = \text{MSE}(\hat{p}_t, p_t). \quad (7)$$

C. World Knowledge Expert and 4D–2D Alignment

Motivation. Predicting future video alone does not guarantee consistency with body motion. We introduce a compact




Fig. 2: **Attention structure.** Within the coupled 4D / World Knowledge block, tokens attend causally within their stream and to causal history in the partner stream. VLM prefix tokens attend only within the prefix; the Action Expert attends to the prefix and both expert streams.

World Knowledge Expert that learns **implicit 2D-foresight tokens** $Z_t^{\text{wk}} \in \mathbb{R}^{M \times d}$, a low-dimensional summary of what the camera should see, **aligned** with the 4D Future Expert’s predictions.

Unified Bidirectional Cross-Attention (UBCA). The 4D and World Knowledge experts are two transformer modules, each carrying temporal tokens. At layer ℓ and timestep i , for module $A \in \{4\text{D}, \text{wk}\}$ with companion B :

$$\tilde{h}_{\ell,i}^A = \text{Attn}(q(h_{\ell,i}^A), K_{\leq i}^A, V_{\leq i}^A) + \text{Attn}(q(h_{\ell,i}^A), K_{\leq i}^B, V_{\leq i}^B), \quad (8)$$

with causal masking over timesteps: each token attends causally within its stream and to causal history in the partner stream. Fig. 2 shows the structured attention pattern.

The world-knowledge pathway produces

$$Z_t^{\text{wk}} = f_{\text{wk}}(o_t, \ell_t, q_t, Z_t^{4\text{D}}). \quad (9)$$

During training, Z_t^{wk} conditions a **frozen** video diffusion transformer g_{wm} . With VAE latents x_0 , noise ϵ , and $\tau \sim \mathcal{U}[0, 1]$:

$$\mathcal{L}_{\text{vid}} = \mathbb{E}_{\tau, \epsilon} \left[\|g_{\text{wm}}(x_\tau, h(Z_t^{\text{wk}})) - (\epsilon - x_0)\|^2 \right]. \quad (10)$$

Only f_{wk} , UBCA parameters, and projection h are updated; g_{wm} is **frozen**.

D. Action Expert

The Action Expert generates action chunk A_t via flow matching, conditioned on the VLM prefix, $Z_t^{4\text{D}}$, Z_t^{wk} , and

q_t . With ground-truth A_t^0 , noise ϵ , and $\tau \sim \mathcal{U}[0, 1]$:

$$\mathcal{L}_{\text{act}} = \mathbb{E}_{\tau, \epsilon} \left[\|v_\theta(A_t^\tau, o_t, \ell_t, q_t, Z_t^{4\text{D}}, Z_t^{\text{wk}}) - (\epsilon - A_t^0)\|^2 \right]. \quad (11)$$

E. Multi-Task Training Objective

$$\mathcal{L} = \lambda_{\text{cur}} \mathcal{L}_{\text{cur}} + \lambda_{4\text{D}} \mathcal{L}_{4\text{D}} + \lambda_{\text{vid}} \mathcal{L}_{\text{vid}} + \lambda_{\text{act}} \mathcal{L}_{\text{act}} + \lambda_{\text{vqa}} \mathcal{L}_{\text{vqa}}. \quad (12)$$

Post-training runs for **10,000** steps. The video world model g_{wm} is **frozen**. All other trainable modules (VLM, 4D Future Expert, World Knowledge Expert, UBCA, projection h , and Action Expert) are updated jointly from step 0. Let η denote the base learning rate for experts and the action head; the VLM uses $\eta/5$.

F. Inference

At deployment: (1) encode (o_t, ℓ_t) with the VLM; (2) run both experts with UBCA to obtain $Z_t^{4\text{D}}$, Z_t^{wk} ; (3) run the Action Expert to output A_t ; (4) **do not** invoke g_{wm} . The extra cost vs. a plain 2D VLA is two lightweight expert forward passes and FK on q_t .

IV. EXPERIMENTS

We evaluate 4DW-VLA on RoboTwin 2.0 [4] against two **representative baseline families** that isolate different gaps in current VLA design. We ask: (1) Does explicit **4D–2D alignment** improve in-domain (Easy) success over a **2D-foresight** VLA that predicts future video but not metric 3D/4D body motion? (2) Does the same recipe improve robustness on the Hard (domain-randomized, OOD) split? (3) Do gains persist across stacking, tool-use scanning, and placement tasks?

A. Experimental Setup

Benchmark. RoboTwin 2.0 is a dual-arm manipulation suite on SAPIEN with domain randomization over clutter, lighting, background texture, table height, and language [4]. Easy = demo_clean; Hard = demo_randomized (OOD). Evaluations use the Aloha-AgileX bimanual embodiment; success rate is the metric.

Tasks. `stack_bowls_three` (geometrically precise stacking), `scan_object` (dual-arm tool use), and `place_bread_skillet` (distinct placement objective).

Baselines. We compare against two published VLAs that represent complementary design points; neither serves as an architectural parent of 4DW-VLA.

InternVLA-A1.5 [9] represents the **2D-foresight VLA** family: temporal reasoning and future **2D video** distillation, but **no explicit 3D/4D body motion** or 4D–2D alignment. We re-implement and train it under the same post-training budget as 4DW-VLA.

LingBot-VLA 2.0 [23] represents the **depth-plus-video VLA** family: **3D (depth)** cues and temporal video generation, but **no coupling** of 3D/4D body trajectories with 2D temporal evolution. We do not reproduce LingBot-VLA 2.0;

TABLE I: **Post-training success rates (%) on RoboTwin 2.0.** Easy = demo_clean; Hard = demo_randomized. Bold = best per column.

Method	stack_bowls		scan_obj.		place_bread	
	E	H	E	H	E	H
InternVLA-A1.5	71	56	45	28	31	20
LingBot-VLA 2.0	78	22	n/a	n/a	n/a	n/a
4DW-VLA	81	58	45	33	39	29

TABLE II: **Shared post-training recipe.**

Item	Setting
Benchmark	RoboTwin 2.0
Post-training budget	10,000 steps
Video world model g_{wm}	Frozen
VLM learning rate	$\eta/5$ (experts / action at η)
Easy / Hard	demo_clean / demo_randomized




Fig. 3: **RoboTwin 2.0 success rates.** Grouped bars for Easy and Hard on three tasks. LingBot-VLA 2.0: n/a on two tasks.

numbers are **as-reported**. Not evaluated on scan_object or place_bread_skillet (n/a).

Training recipe. 4DW-VLA and other methods share a **10,000-step** post-training budget. The video world model g_{wm} is **frozen**. Expert and action-head parameters use base learning rate η ; the VLM uses $\eta/5$ (Table II).

B. Main Results

Fig. 3 visualizes Table I.

stack_bowls_three. 4DW-VLA reaches 81%/58% (Easy/Hard) vs. 71%/56% and 78%/22% for other methods. On Easy all three are competitive; on Hard, 4DW-VLA holds 58% while one baseline collapses to 22%, consistent with **4D-2D aligned body foresight** under visual randomization.

scan_object. 4DW-VLA ties on Easy (45%) and leads on Hard (33% vs. 28%).

place_bread_skillet. 4DW-VLA reaches 39%/29% vs. 31%/20% for other methods (+8/+9 points). LingBot-VLA 2.0 has no reported result on this task.

4DW-VLA achieves the highest success rate in every evaluated column. Hard gains over the best baseline per task are +2, +5, and +9 points, respectively.

C. Discussion

Easy scores match or exceed other methods on all three tasks, indicating that 4D and aligned 2D auxiliary heads need not hurt in-domain learning. Consistent Hard improvements support the hypothesis that **explicit 4D-2D coupling** matters more under domain shift than **2D video foresight alone**. LingBot-VLA 2.0 numbers are externally reported values.

V. LIMITATIONS

- **Body-centric 4D, not full scene geometry.** FK trajectories capture the robot’s own motion, not independent object dynamics or extrinsic scene structure.
- **FK accuracy depends on calibration.** URDF-hardware mismatch propagates into 4D labels; sim-to-real kinematics gap is unstudied.
- **World-model distillation quality.** Removing g_{wm} at inference assumes Z_t^{wk} captures sufficient visual dynamics.
- **UBCA ablations pending.** Late concatenation and uni-directional cross-attention are not yet compared.
- **Limited task set.** Three RoboTwin 2.0 tasks only; trial counts, seeds, and confidence intervals are not reported.
- **Simulation only.** All Table I numbers are from simulation; real-robot transfer is untested.

VI. CONCLUSION

We presented **4DW-VLA**, a **4D and world-model aligned VLA** for robotic manipulation. Its central component is a **4D Future Expert** that derives lightweight 7D body trajectories from FK and is **retained at deployment**, unlike prior FK-based approaches that discard the 4D branch after training. Coupled with a World Knowledge Expert through bidirectional cross-attention and a frozen video world model used only during training, the architecture produces a small, fast policy with 4D-aligned visual foresight and no deployment-time video generation. On RoboTwin 2.0, 4DW-VLA matches or exceeds a **2D-foresight** baseline and a **depth-plus-video** baseline, with the strongest advantage under domain randomization. Future work will ablate the 4D expert (retained vs. discarded, 7D vs. 3D), extend evaluation, and validate on real hardware.

REFERENCES

- [1] Kevin Black et al. π_0 : A vision-language-action flow model for general robot control. *arXiv preprint arXiv:2410.24164*, 2024.
- [2] Anthony Brohan, Noah Brown, Justice Carbajal, et al. RT-1: Robotics transformer for real-world control at scale. In *Robotics: Science and Systems (RSS)*, 2023.
- [3] Anthony Brohan, Noah Brown, Justice Carbajal, et al. RT-2: Vision-language-action models transfer web knowledge to robotic control. In *Conference on Robot Learning (CoRL)*, 2023.
- [4] Tianyou Chen et al. RoboTwin 2.0: A scalable data generator and benchmark with strong domain randomization for robust bimanual robotic manipulation. *arXiv preprint arXiv:2506.18088*, 2025.
- [5] Cheng Chi, Siyuan Feng, Yilun Du, Zhenjia Xu, Eric Cousineau, Benjamin Burchfiel, and Shuran Song. Diffusion policy: Visuomotor policy learning via action diffusion. In *Robotics: Science and Systems (RSS)*, 2023.
- [6] Theophile Gervet et al. Act3D: 3D feature field transformers for multi-task robotic manipulation. In *Conference on Robot Learning (CoRL)*, 2023.
- [7] David Ha and Jürgen Schmidhuber. World models. *arXiv preprint arXiv:1803.10122*, 2018.

- [8] Nicklas Hansen, Hao Su, and Xiaolong Wang. TD-MPC2: Scalable, robust world models for continuous control. In *International Conference on Learning Representations (ICLR)*, 2024.
- [9] InternVLA-A1.5 Team. InternVLA-A1.5: Unifying understanding, latent foresight, and action for compositional generalization. <https://internrobotics.github.io/internvla-a15.github.io/>, 2026.
- [10] Alexander Khazatsky et al. DROID: A large-scale in-the-wild robot manipulation dataset. In *Robotics: Science and Systems (RSS)*, 2024.
- [11] Moo Jin Kim et al. OpenVLA: An open-source vision-language-action model. In *Conference on Robot Learning (CoRL)*, 2024.
- [12] Jonathan Levy et al. Simulation distillation: Pretraining world models in simulation for rapid real-world adaptation. *arXiv preprint arXiv:2603.15759*, 2026.
- [13] Songming Liu et al. RDT-1B: A diffusion foundation model for bimanual manipulation. *arXiv preprint arXiv:2410.07864*, 2024.
- [14] NVIDIA. Pretrained to imagine, fine-tuned to act: The rise of world-action models. NVIDIA Developer Blog, 2026.
- [15] others. SpatialVLA: Exploring spatial representations for vision-language-action model. *arXiv preprint arXiv:2501.15830*, 2025.
- [16] others. Dreaming when necessary: Advancing world action models with adaptive multi-modal reasoning. *arXiv preprint arXiv:2606.07089*, 2026.
- [17] others. 4D-WAM: Infusing spatiotemporal awareness into world action models through trajectory fields. *arXiv preprint arXiv:2608.08023*, 2026.
- [18] others. LaWAM: Latent world action models for efficient dynamics-aware robot policies. *arXiv preprint arXiv:2606.15768*, 2026.
- [19] others. Lift3D-VLA: Lifting VLA models to 3D geometry and dynamics-aware manipulation. *arXiv preprint arXiv:2607.06564*, 2026.
- [20] others. Pri4R: Learning world dynamics for vision-language-action models with privileged 4D representation. *arXiv preprint arXiv:2603.01549*, 2026.
- [21] others. World action models (WAM): A survey. <https://world-action-models.github.io/>, 2026.
- [22] Jingjing Qian et al. 4D-VLA: Spatiotemporal vision-language-action pretraining with cross-scene calibration. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2025.
- [23] Robbyant. LingBot-VLA 2.0: From foundation to application. <https://github.com/Robbyant/lingbot-vla-v2>, 2026.
- [24] Quan Vuong et al. Open X-Embodiment: Robotic learning datasets and RT-X models, 2023.
- [25] Homer Rich Walke et al. BridgeData V2: A dataset for robot learning at scale. In *Conference on Robot Learning (CoRL)*, 2023.
- [26] Rui Yang, Guojian Chen, Chuan Wen, and Yihua Gao. FP3: A 3D foundation policy for robotic manipulation. *arXiv preprint arXiv:2503.08950*, 2025. ICRA 2026 Best Paper.
- [27] Yanjie Ze et al. 3D diffusion policy: Generalizable visuomotor policy learning via simple 3D representations. In *Robotics: Science and Systems (RSS)*, 2024.
- [28] Haoyu Zhao, Xingyue Zhao, Siteng Huang, et al. RynnWorld-4D: 4D embodied world models for robotic manipulation. *arXiv preprint arXiv:2607.06559*, 2026.
- [29] Tony Z. Zhao, Vikash Kumar, Sergey Levine, and Chelsea Finn. Learning fine-grained bimanual manipulation with low-cost hardware. In *Robotics: Science and Systems (RSS)*, 2023.